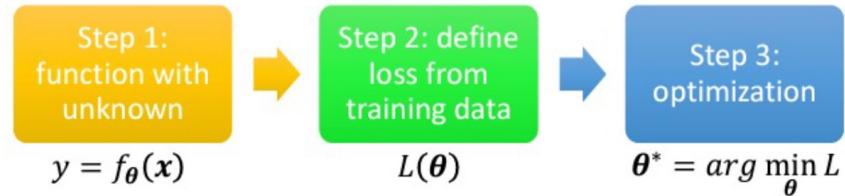


LTTI.00.032 – Machine Learning in Synthetic Biology

Recap

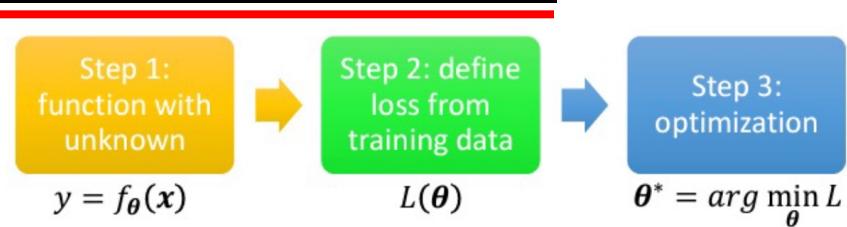
- The essence of deep learning is to find best weights for the network that minimizes the error (loss)

Recap



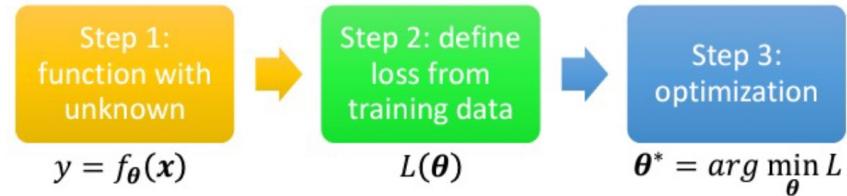
- The essence of deep learning is to find best weights for the network that minimizes the error (loss)

Recap



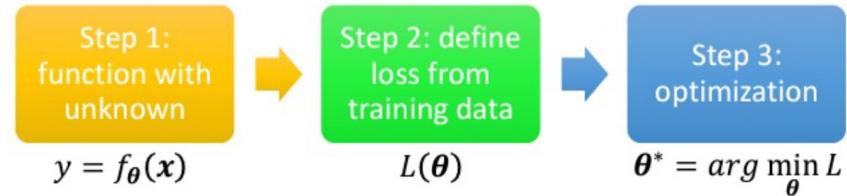
- The essence of deep learning is to find best weights for the network that minimizes the error (loss)
- This is done via an iterative process where weights are updated in each iteration in a direction that minimizes the loss

Recap



- The essence of deep learning is to find best weights for the network that minimizes the error (loss)
- This is done via an iterative process where weights are updated in each iteration in a direction that minimizes the loss
- To find this direction we need the slope of the loss function for a given weight. This is achieved by computing the derivatives (gradient)

Recap



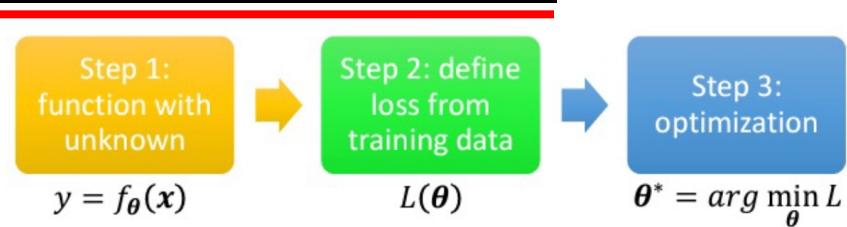
- The essence of deep learning is to find best weights for the network that minimizes the error (loss)
- This is done via an iterative process where weights are updated in each iteration in a direction that minimizes the loss
- To find this direction we need the slope of the loss function for a given weight. This is achieved by computing the derivatives (gradient)
- It is computationally expensive to compute derivatives for millions of weights

Recap

- Backpropagation makes this computation possible by using the chain rule in calculus

Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *nature*, 323(6088), 533-536.

Recap

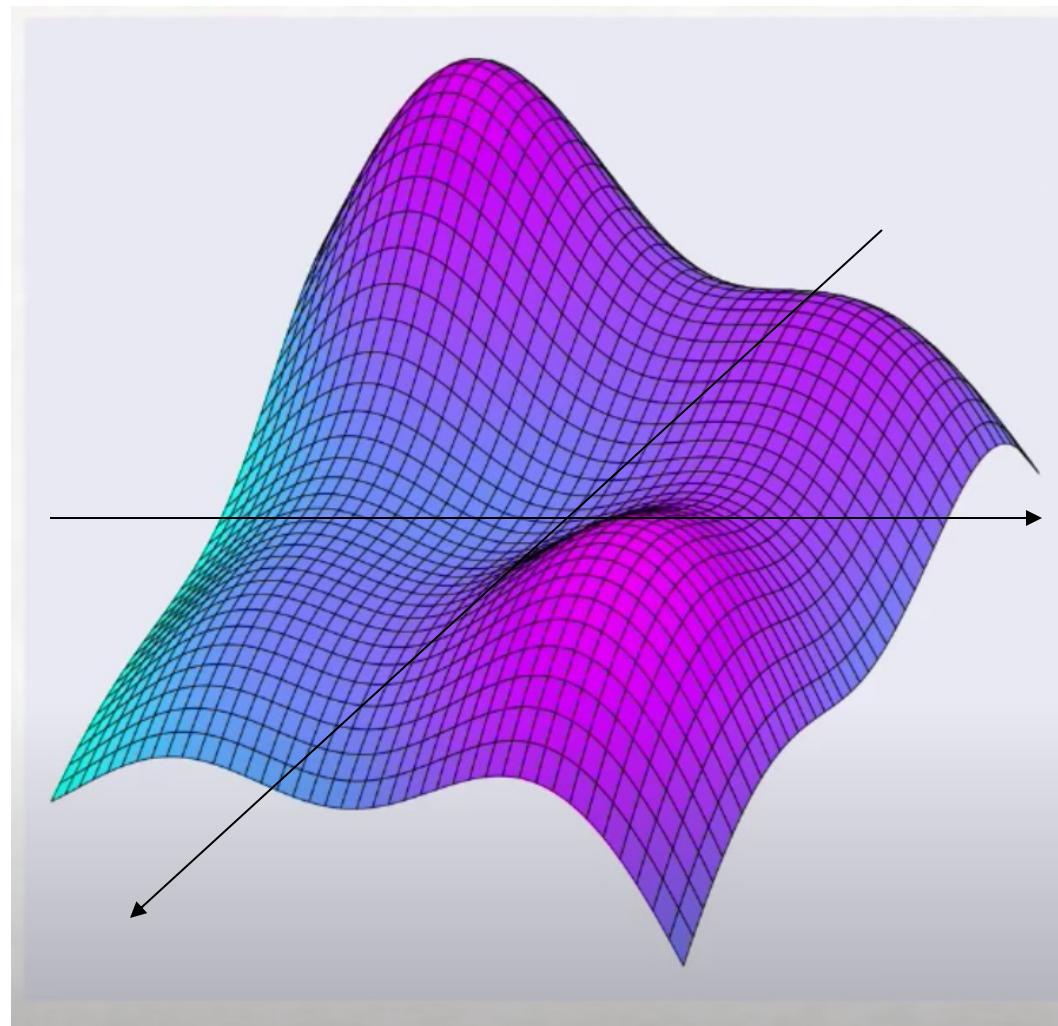


- Backpropagation makes this computation possible by using the chain rule in calculus

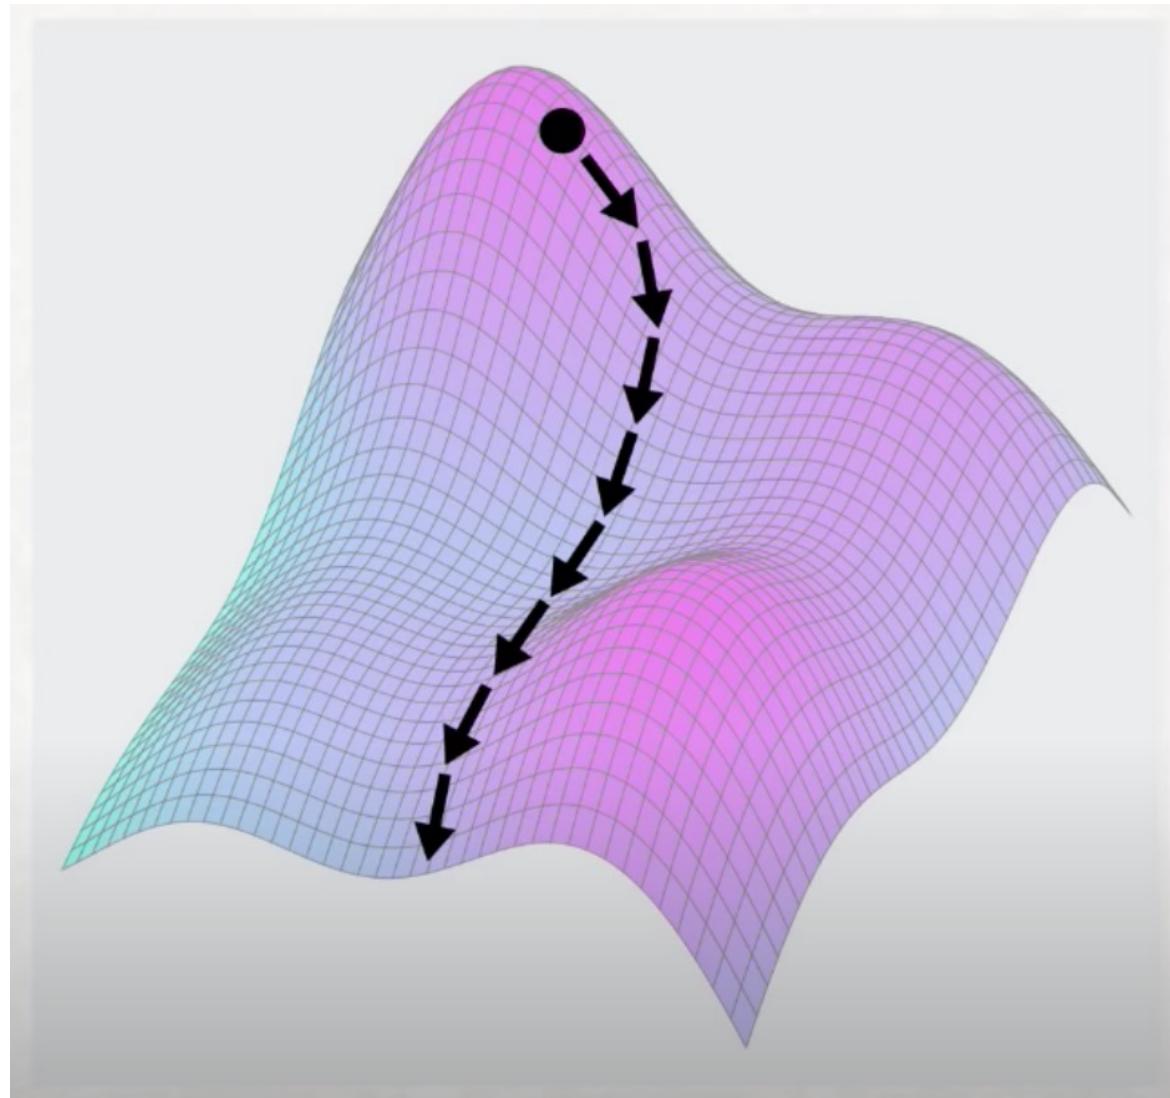
Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *nature*, 323(6088), 533-536.

- Gradient decent is a general term for calculating the gradient and updating the weights

Gradient Decent

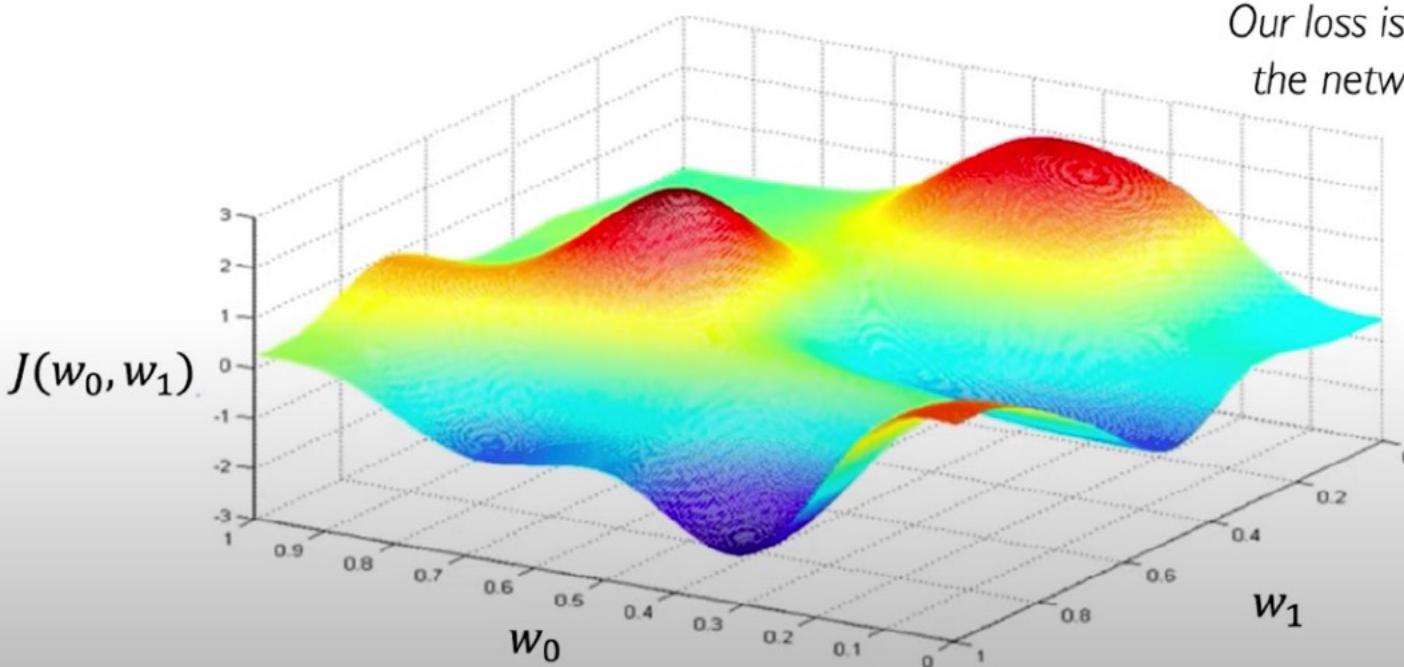


Gradient Decent



Loss optimization

$$\mathbf{W}^* = \underset{\mathbf{W}}{\operatorname{argmin}} J(\mathbf{W})$$

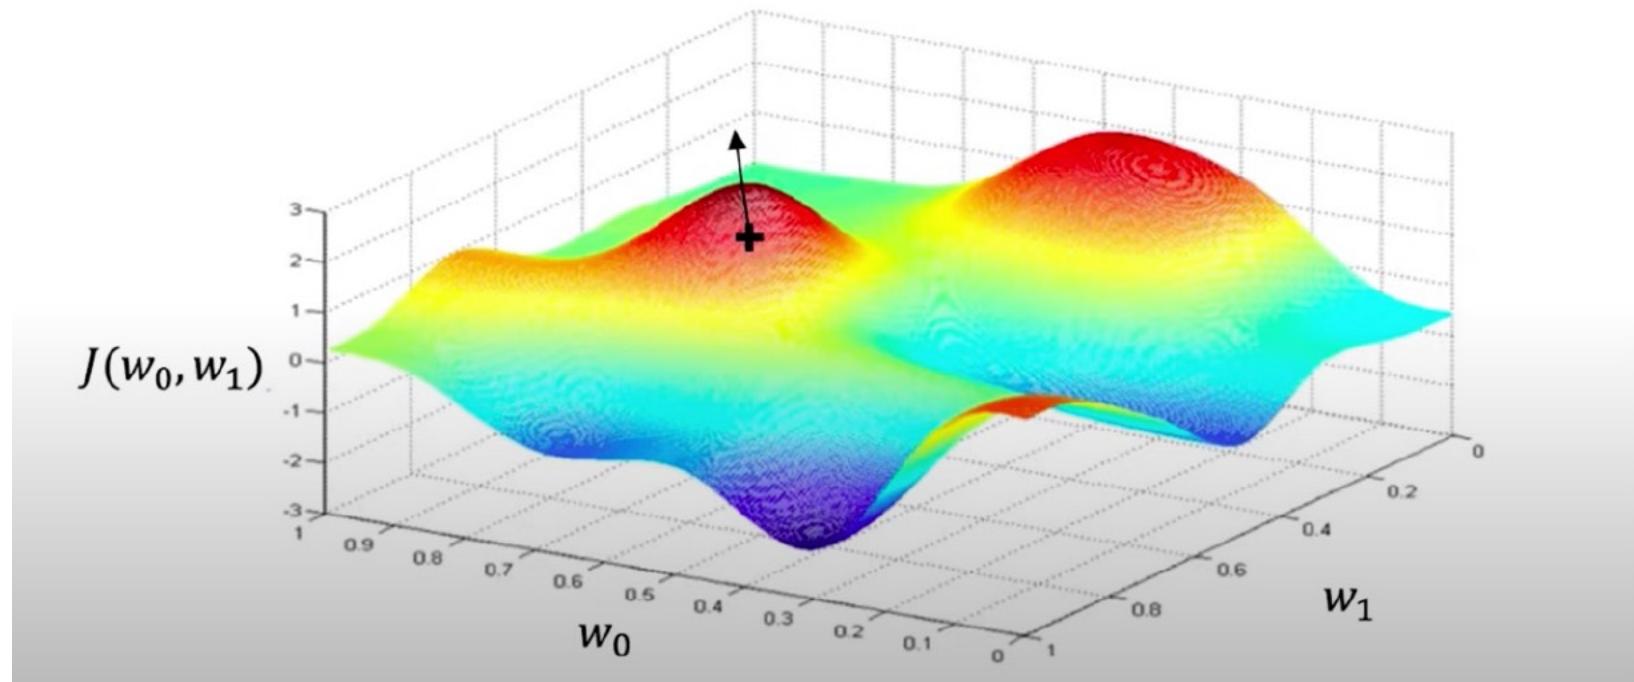


Remember:

*Our loss is a function of
the network weights!*

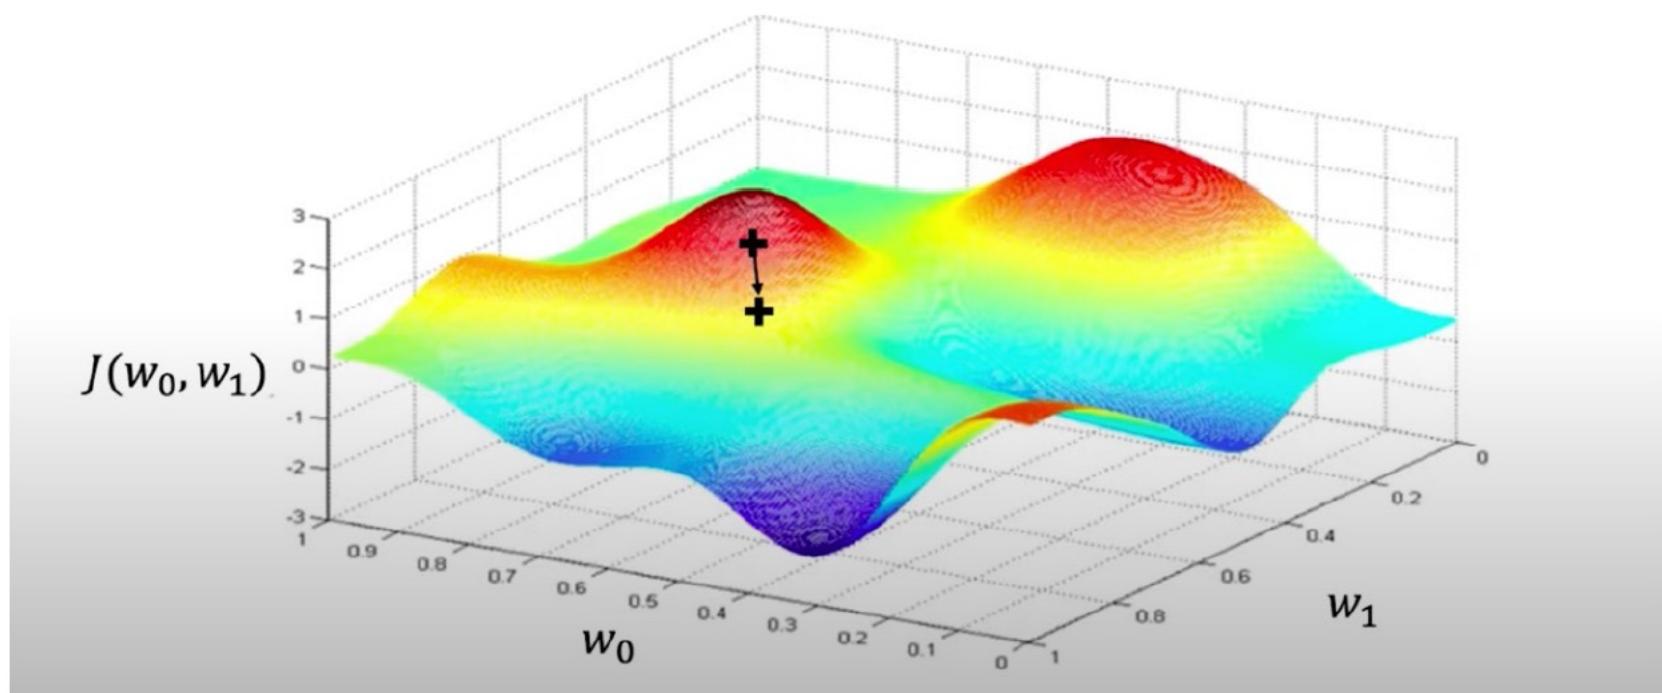
Loss optimization

Compute gradient, $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$



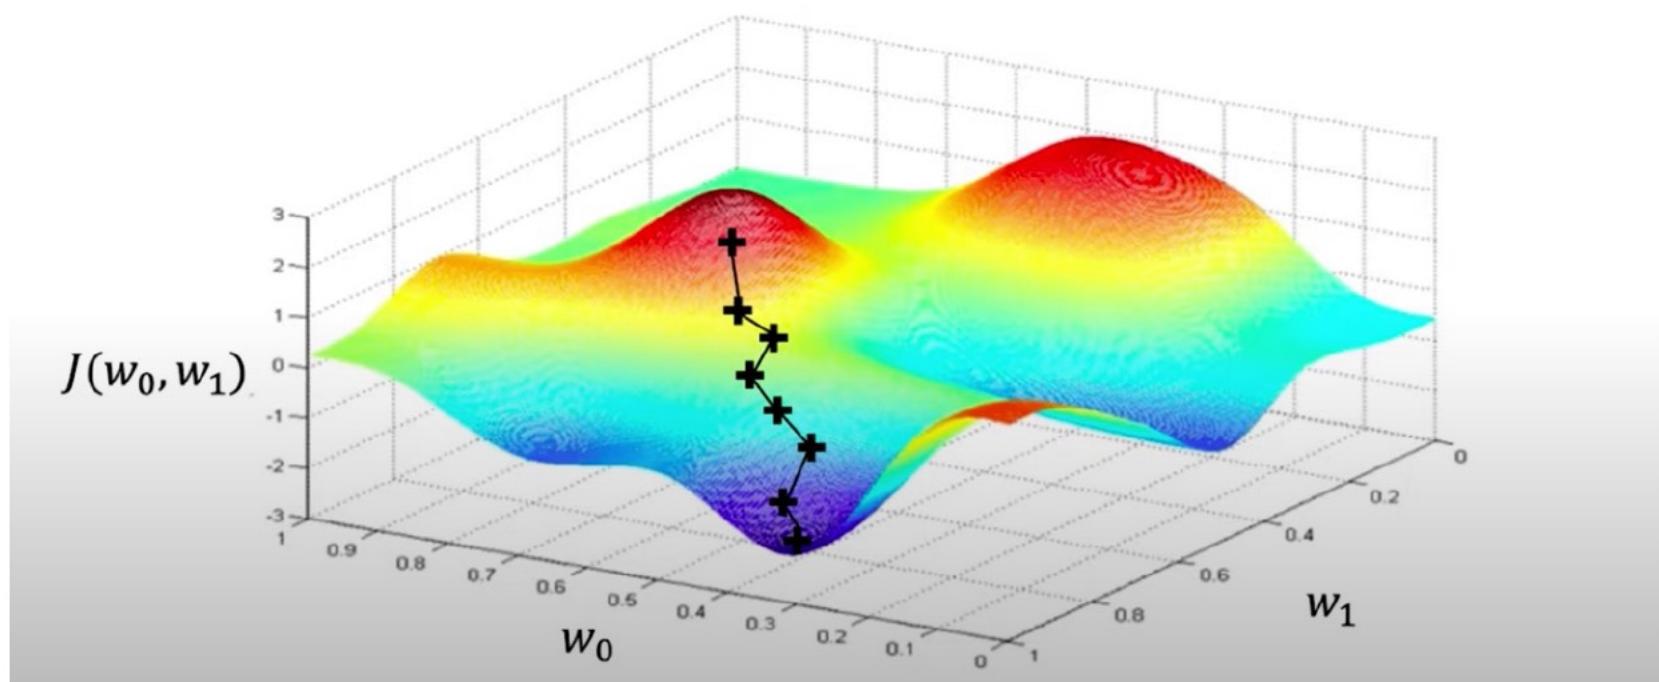
Loss optimization

Take small step in opposite direction of gradient



Gradient decent

Repeat until convergence

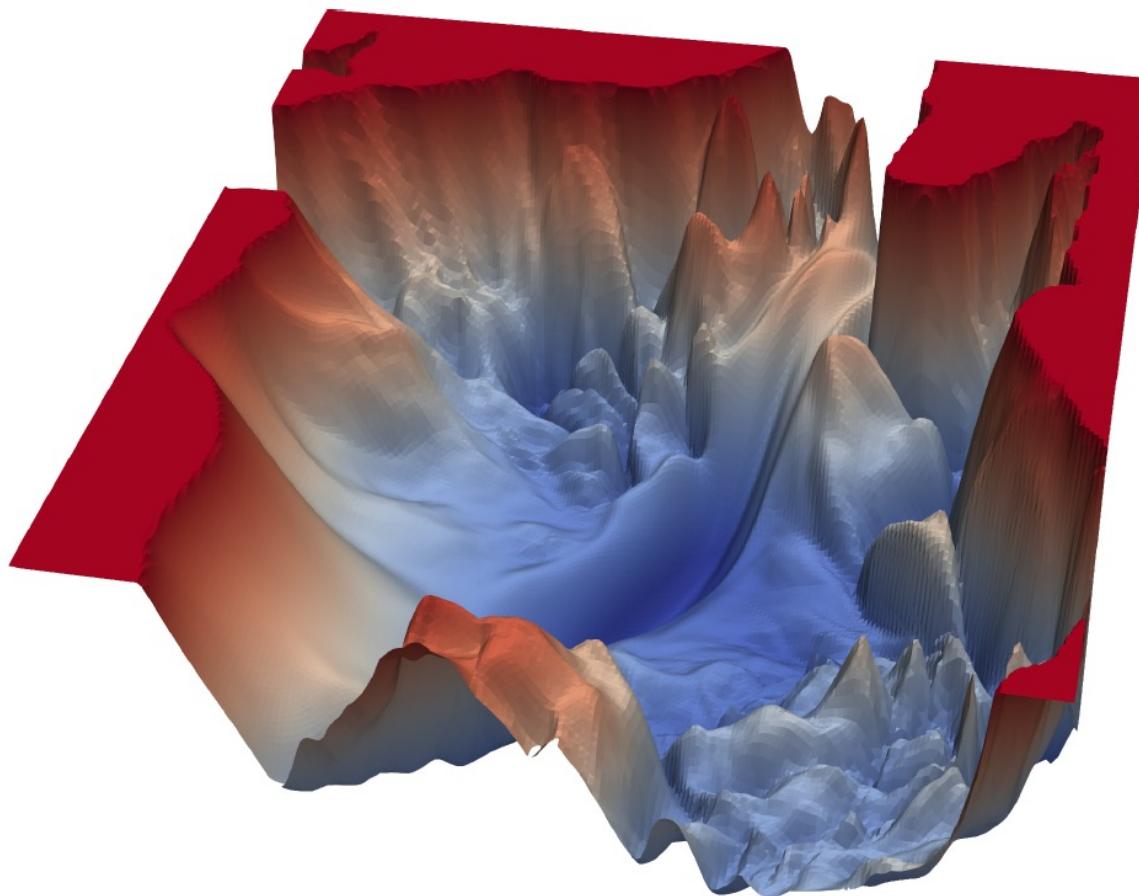


Gradient decent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{w}}$
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{w}}$
5. Return weights

Gradient decent



Loss function optimization

Remember:

Optimization through gradient descent

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$

Loss function optimization

Remember:

Optimization through gradient descent

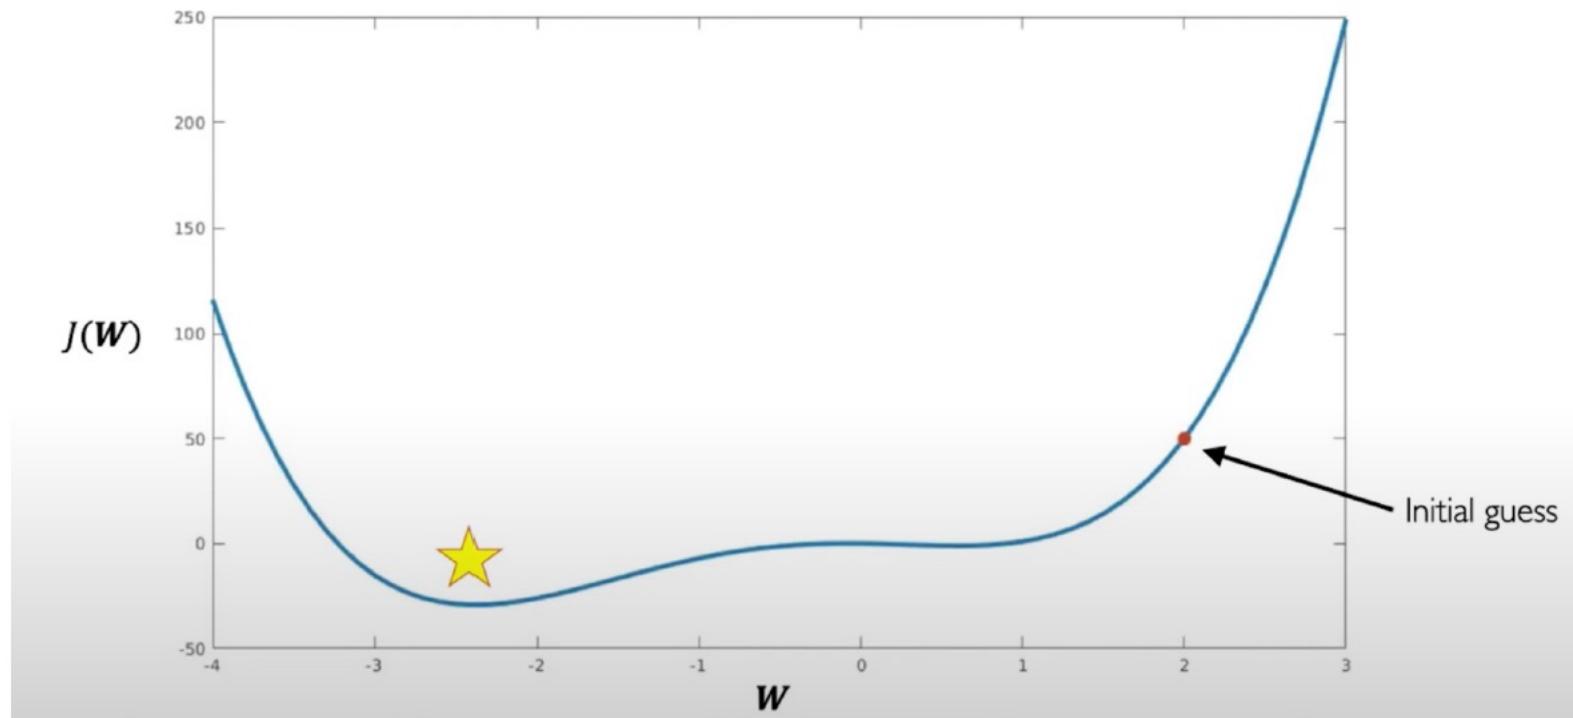
$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$$



How can we set the learning rate?

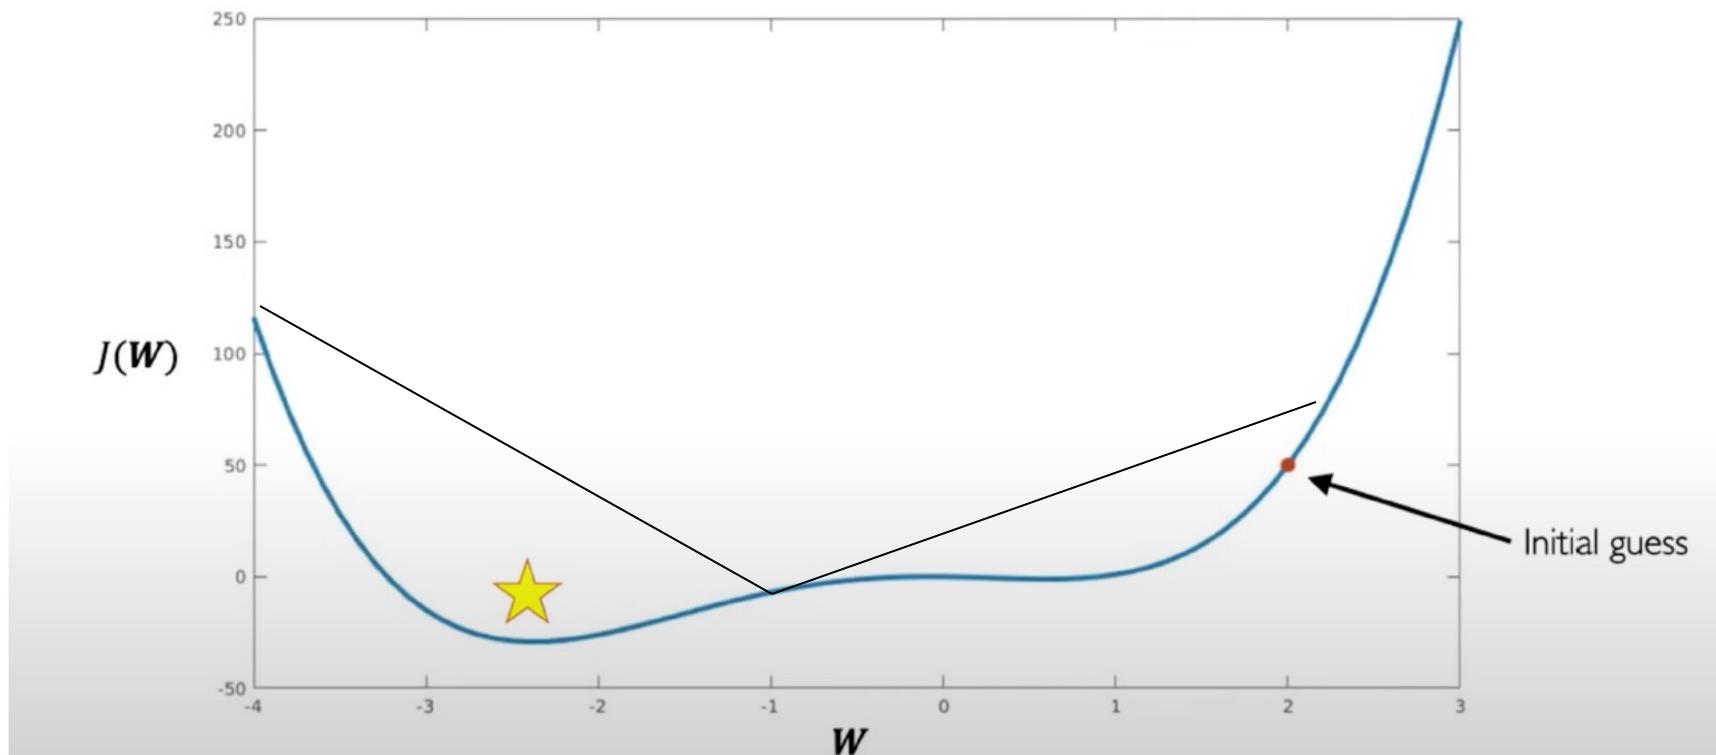
Setting the learning rate

Small learning rate converges slowly and gets stuck in false local minima



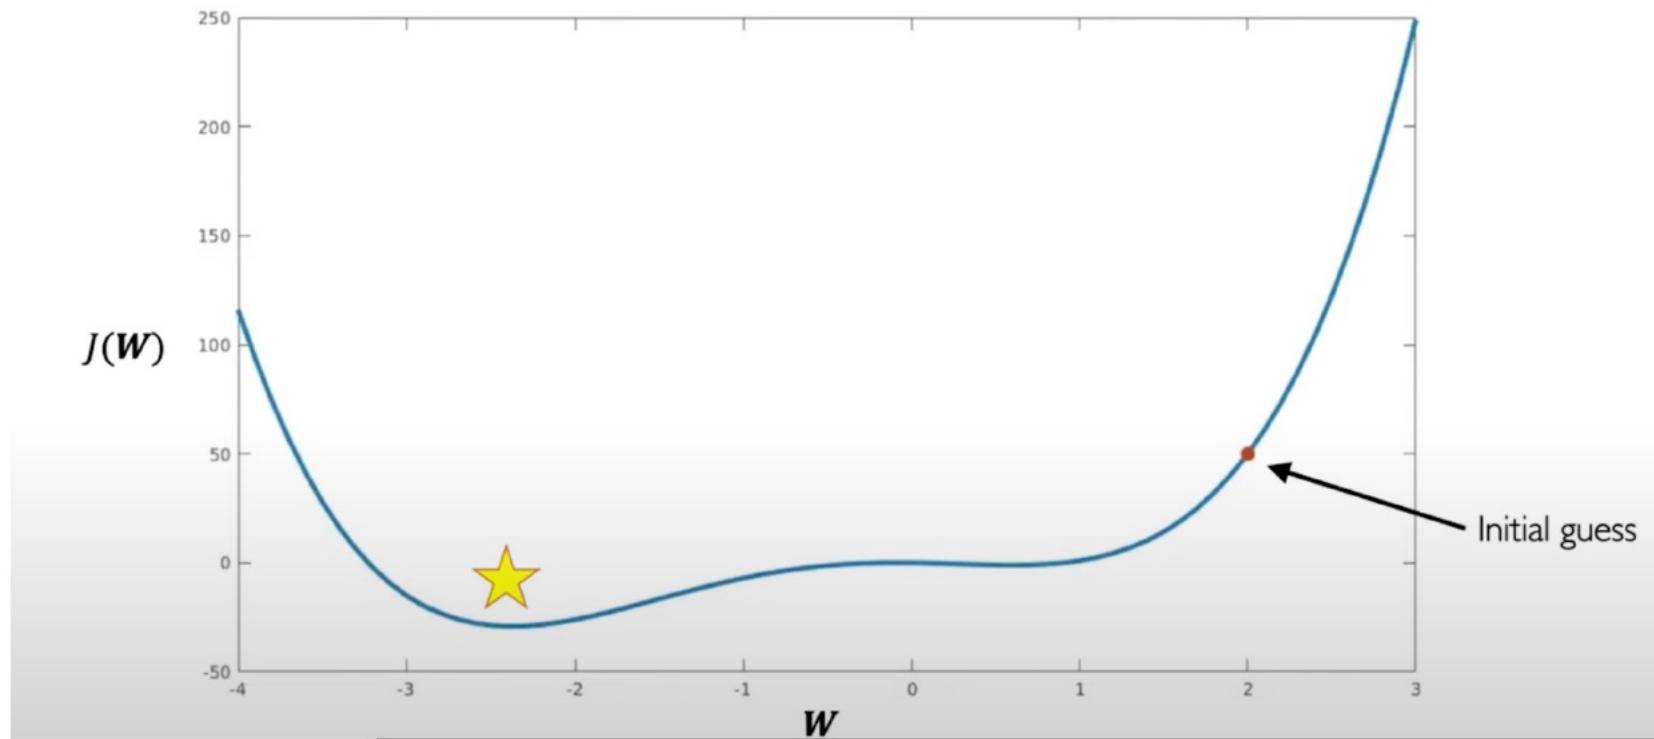
Setting the learning rate

Large learning rates overshoot, become unstable and diverge



Setting the learning rate

Stable learning rates converge smoothly and avoid local minima



How to deal with this

Idea 1.

Try lots of different learning rates and see what works “just right”

How to deal with this

Idea 1.

Try lots of different learning rates and see what works “just right”

Idea 2.

Design an adaptive learning rate that adapts to the landscape

Gradient decent algorithms

SGD

Adam

Adadelta

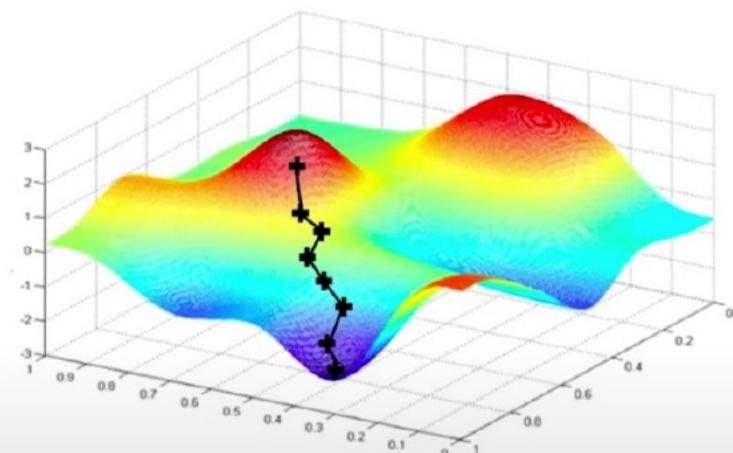
Adagrad

RMSProp

Gradient decent

Algorithm

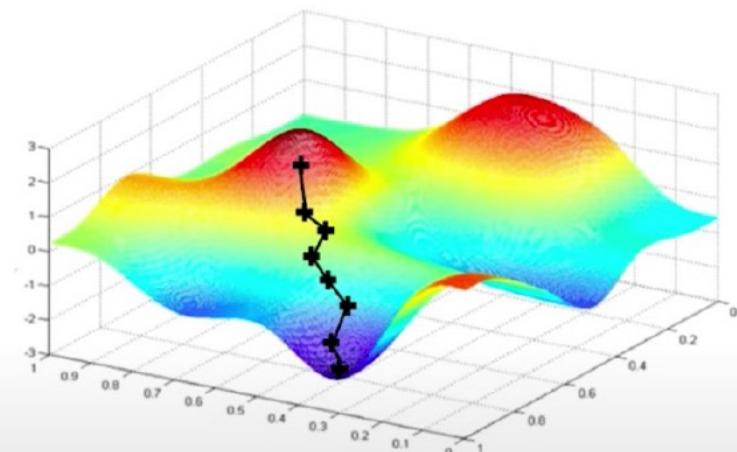
1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights



Gradient decent

Algorithm

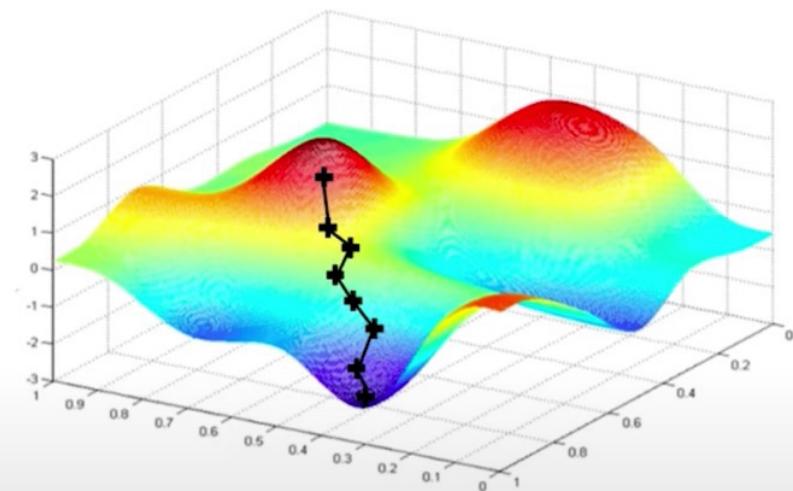
1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights



Stochastic Gradient Decent

Algorithm

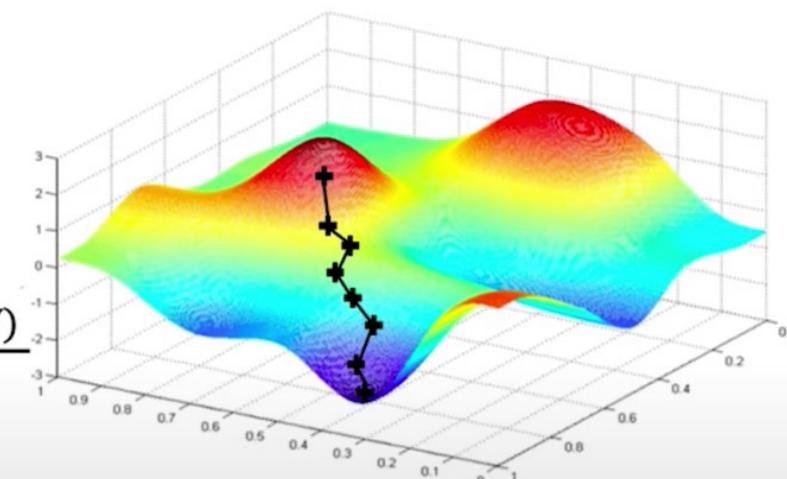
1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick single data point i
4. Compute gradient, $\frac{\partial J_i(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights



Stochastic Gradient Decent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick batch of B data points
4. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}} = \frac{1}{B} \sum_{k=1}^B \frac{\partial J_k(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights



Mini-batches while training

- More accurate estimation of gradient
- Smoother convergence
- Allows for larger learning rates
- Leads to fast training

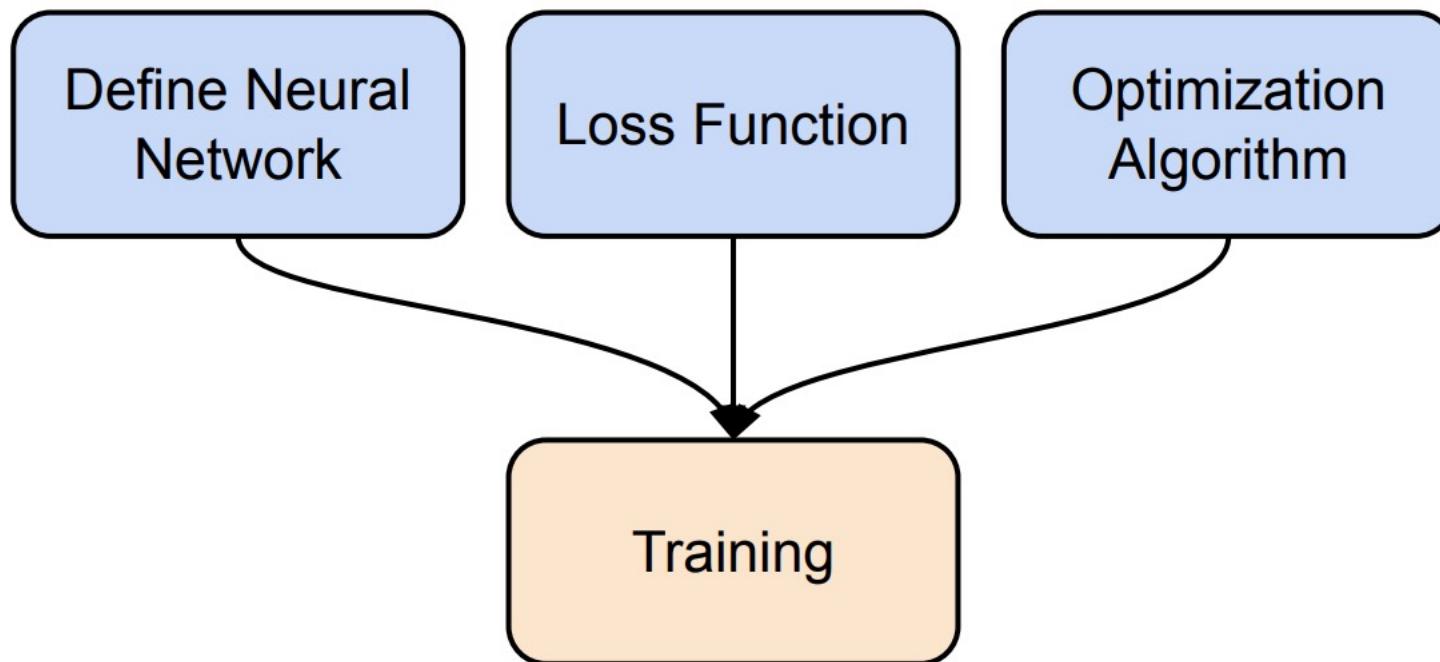
Practical

What is PyTorch?

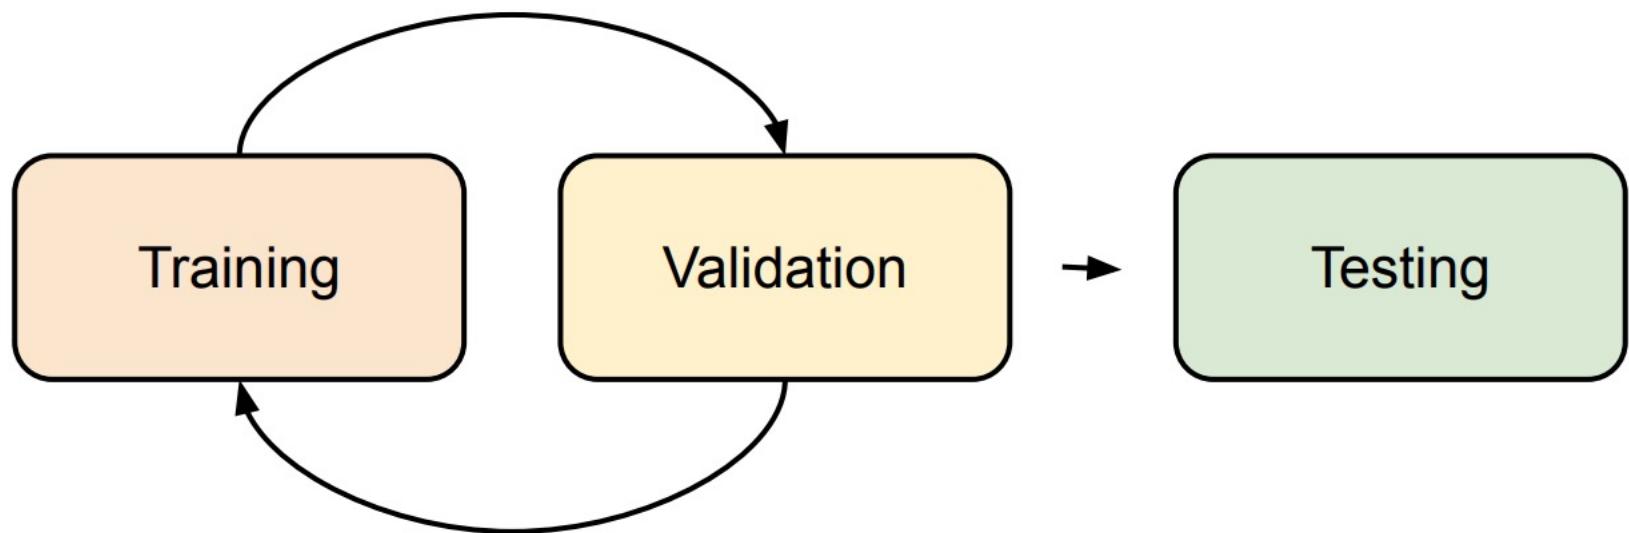
- An **machine learning framework** in Python
- Two main features:
 - N-dimensional **Tensor** computation on **GPUs**
 - Automatic differentiation for training deep neural networks



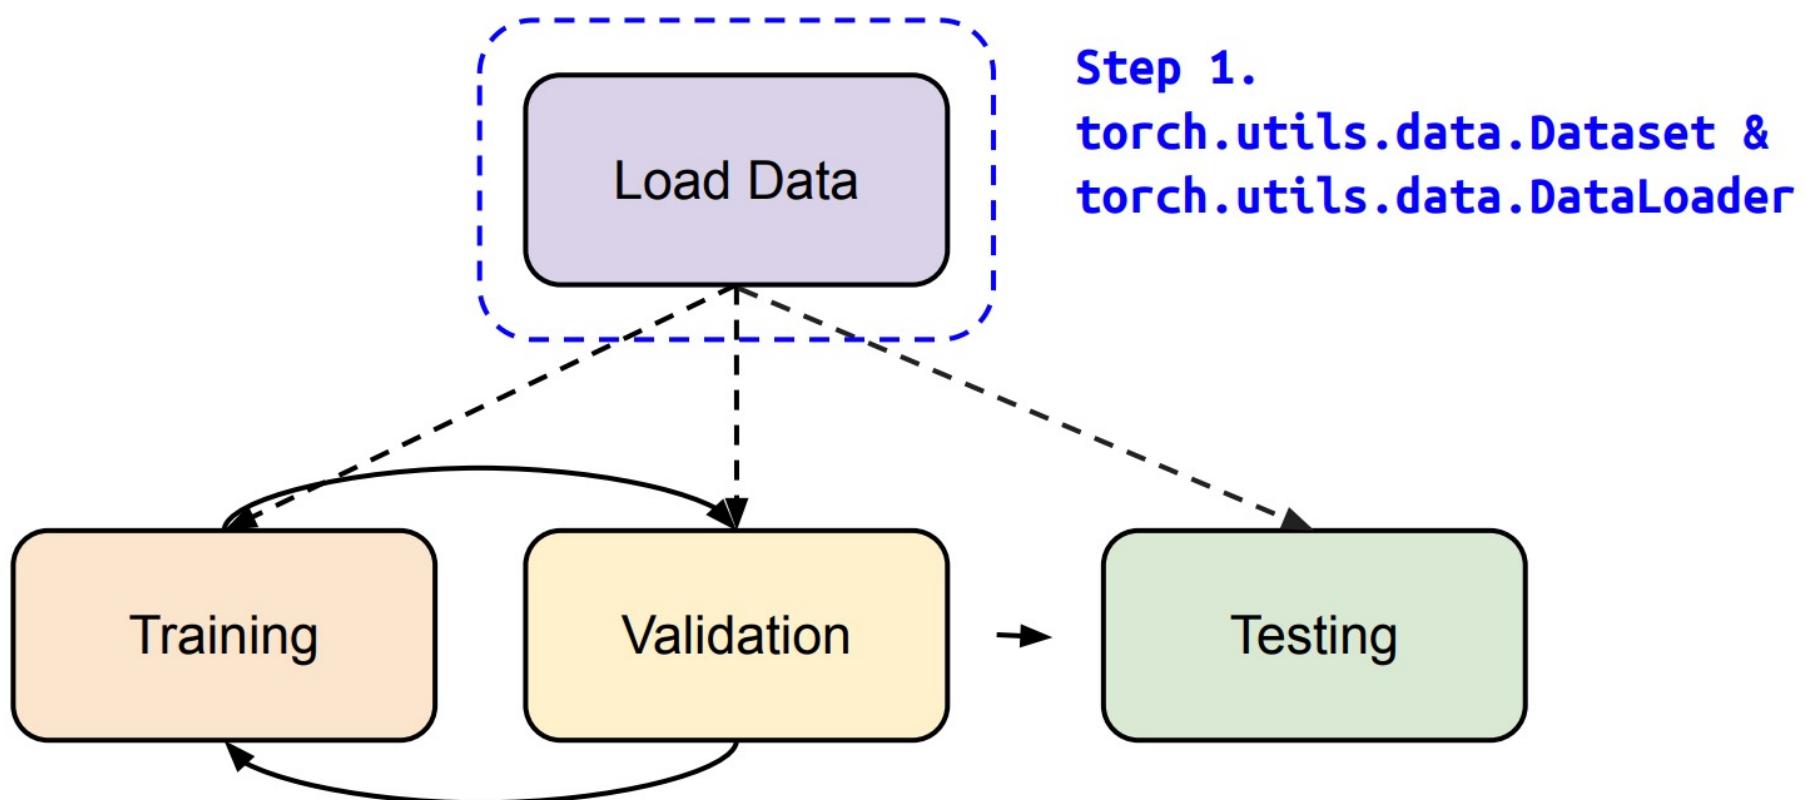
Training Neural Networks



Training & Testing Neural Networks



Training & Testing Neural Networks in PyTorch



Dataset & Dataloader

- Dataset: stores data samples and expected values
- Dataloader: groups data in batches, enables multiprocessing
- `dataset = MyDataset(file)`
- `dataloader = DataLoader(dataset, batch_size, shuffle=True)`



Training: True
Testing: False

Dataset & Dataloader

- Dataset: stores data samples and expected values
- Dataloader: groups data in batches, enables multiprocessing
- `dataset = MyDataset(file)`
- `dataloader = DataLoader(dataset, batch_size, shuffle=True)`



Training: True
Testing: False

Dataset & Dataloader

```
from torch.utils.data import Dataset, DataLoader
```

```
class MyDataset(Dataset):  
    def __init__(self, file):  
        self.data = ...  
  
    def __getitem__(self, index):  
        return self.data[index]  
  
    def __len__(self):  
        return len(self.data)
```



Read data & preprocess



Returns one sample at a time



Returns the size of the dataset

Dataset & Dataloader

```
from torch.utils.data import Dataset, DataLoader
```

```
class MyDataset(Dataset):
    def __init__(self, file):
        self.data = ...
    def __getitem__(self, index):
        return self.data[index]
    def __len__(self):
        return len(self.data)
```



Read data & preprocess



Returns one sample at a time

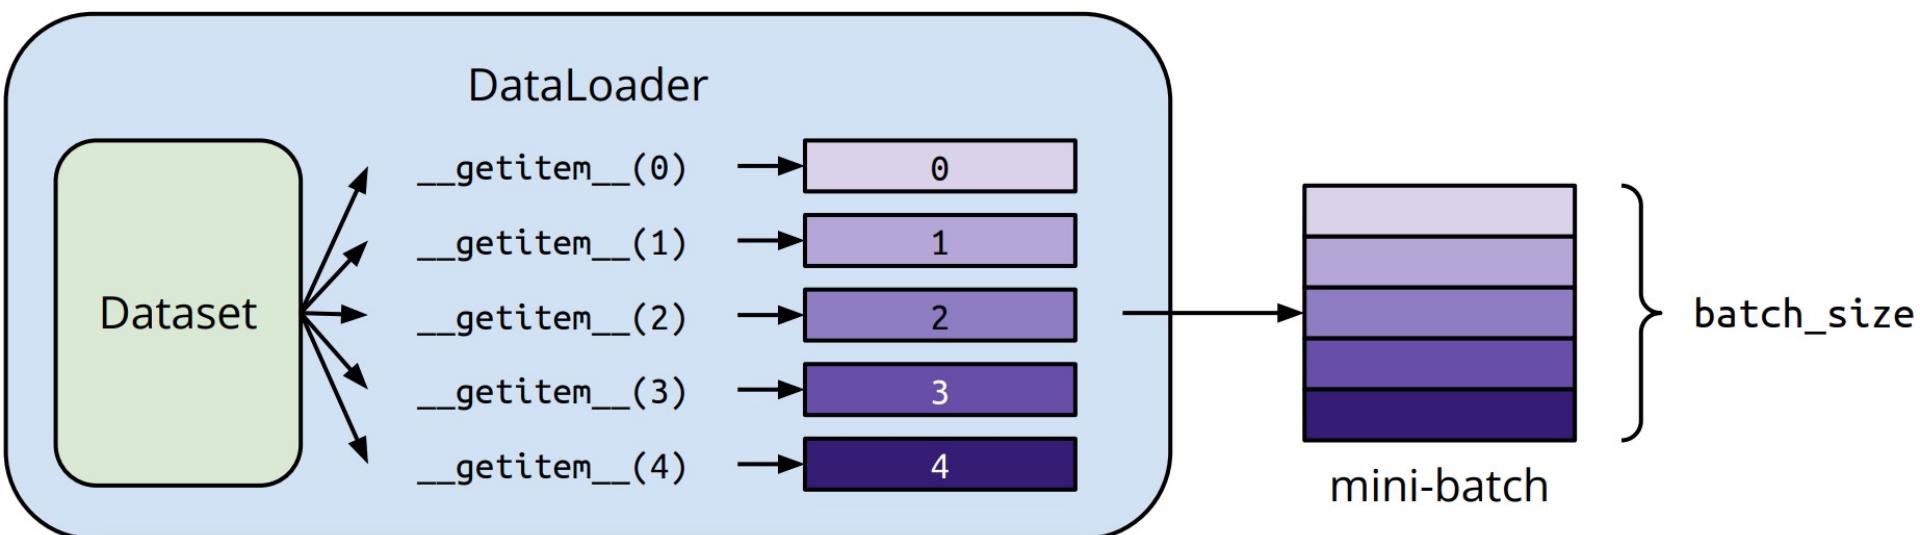


Returns the size of the dataset

Dataset & Dataloader

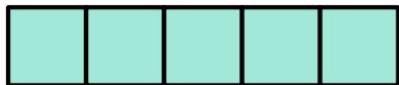
```
dataset = MyDataset(file)
```

```
dataloader = DataLoader(dataset, batch_size=5, shuffle=False)
```

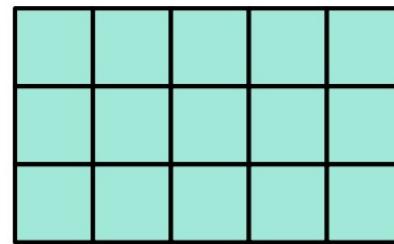


Tensors

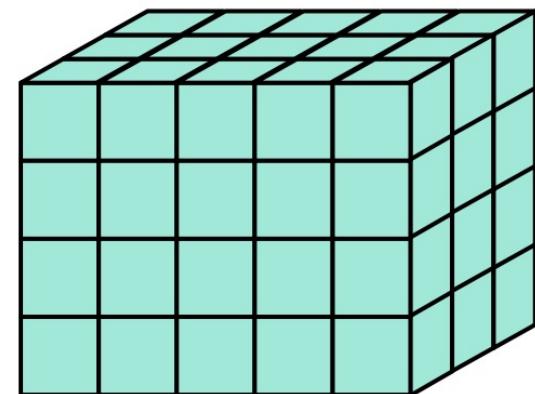
- High-dimensional matrices (arrays)



1-D tensor
e.g. audio



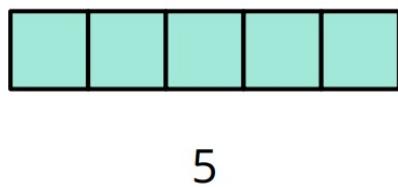
2-D tensor
e.g. black&white
images



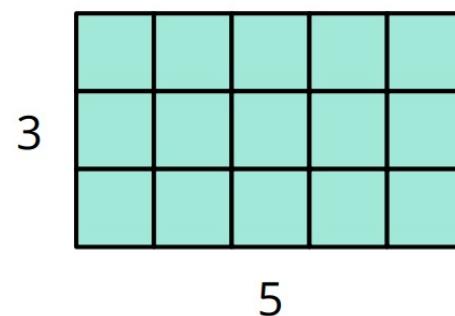
3-D tensor
e.g. RGB images

Tensors – Shape of Tensors

- Check with `.shape()`

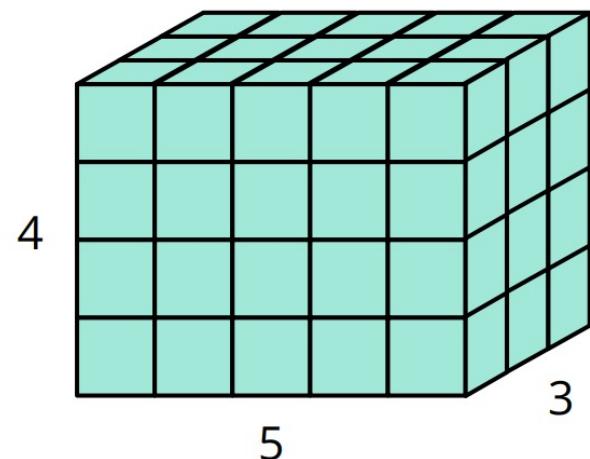


5

 $(5,)$
dim 0

3

5

 $(3, 5)$
dim 0 dim 1

4

5

3

 $(4, 5, 3)$
dim 0 dim 1 dim 2

Note: **dim** in PyTorch == **axis** in NumPy

Tensors – Creating Tensors

- Directly from data (list or numpy.ndarray)

```
x = torch.tensor([[1, -1], [-1, 1]])
```

```
tensor([[1., -1.],  
       [-1., 1.]])
```

```
x = torch.from_numpy(np.array([[1, -1], [-1, 1]]))
```

- Tensor of constant zeros & ones

```
x = torch.zeros([2, 2])
```

```
tensor([[0., 0.],  
       [0., 0.]])
```

```
x = torch.ones([1, 2, 5])
```

shape

```
tensor([[[1., 1., 1., 1., 1.],  
        [1., 1., 1., 1., 1.]]])
```

Tensors – Common Operations

Common arithmetic functions are supported, such as:

- Addition

$$z = x + y$$

- Summation

$$y = x.sum()$$

- Subtraction

$$z = x - y$$

- Mean

$$y = x.mean()$$

- Power

$$y = x.pow(2)$$

Tensors – Common Operations

- **Transpose:** transpose two specified dimensions

```
>>> x = torch.zeros([2, 3])
```

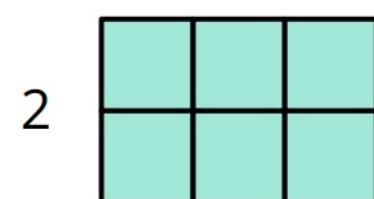
```
>>> x.shape
```

```
torch.Size([2, 3])
```

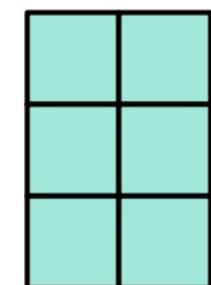
```
>>> x = x.transpose(0, 1)
```

```
>>> x.shape
```

```
torch.Size([3, 2])
```



3
2



2
3

Tensors – Common Operations

- **Squeeze**: remove the specified dimension with length = 1

```
>>> x = torch.zeros([1, 2, 3])
```

```
>>> x.shape
```

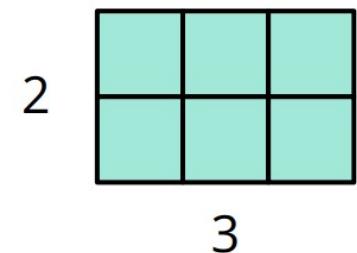
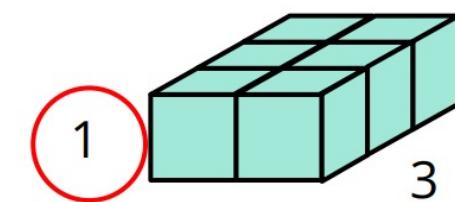
```
torch.Size([1, 2, 3])
```

```
>>> x = x.squeeze(0)
```

(dim = 0)

```
>>> x.shape
```

```
torch.Size([2, 3])
```



Tensors – Common Operations

- **Unsqueeze**: expand a new dimension

```
>>> x = torch.zeros([2, 3])
```

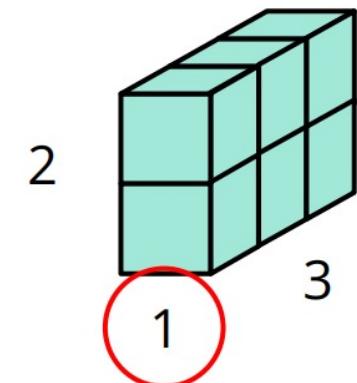
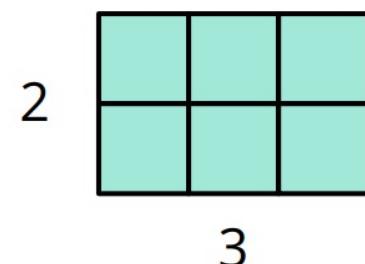
```
>>> x.shape
```

```
torch.Size([2, 3])
```

```
>>> x = x.unsqueeze(1)      (dim = 1)
```

```
>>> x.shape
```

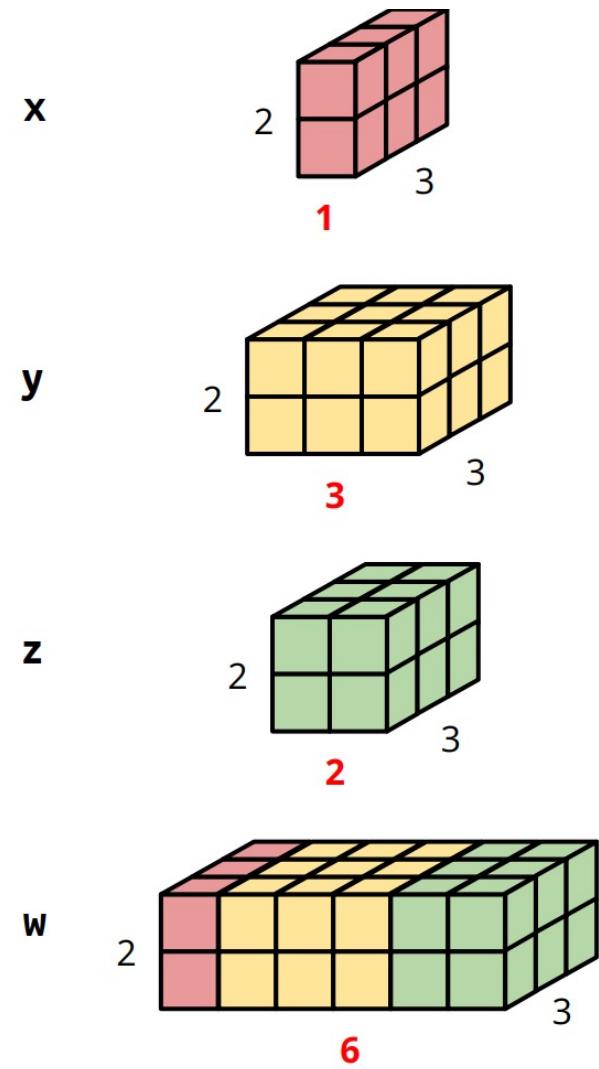
```
torch.Size([2, 1, 3])
```



Tensors – Common Operations

- **Cat:** concatenate multiple tensors

```
>>> x = torch.zeros([2, 1, 3])
>>> y = torch.zeros([2, 3, 3])
>>> z = torch.zeros([2, 2, 3])
>>> w = torch.cat([x, y, z], dim=1)
>>> w.shape
torch.Size([2, 6, 3])
```



more operators: <https://pytorch.org/docs/stable/tensors.html>

Tensors – Data Type

- Using different data types for model and data will cause errors.

Data type	dtype	tensor
32-bit floating point	<code>torch.float</code>	<code>torch.FloatTensor</code>
64-bit integer (signed)	<code>torch.long</code>	<code>torch.LongTensor</code>

Tensors – PyTorch vs Numpy

- Similar attributes

PyTorch	NumPy
<code>x.shape</code>	<code>x.shape</code>
<code>x.dtype</code>	<code>x.dtype</code>

Tensors – PyTorch vs Numpy

- Many functions have the same names as well

PyTorch	NumPy
<code>x.reshape / x.view</code>	<code>x.reshape</code>
<code>x.squeeze()</code>	<code>x.squeeze()</code>
<code>x.unsqueeze(1)</code>	<code>np.expand_dims(x, 1)</code>

Tensors – Device

- Tensors & modules will be computed with **CPU** by default

Use **.to()** to move tensors to appropriate devices.

- CPU

```
x = x.to('cpu')
```

- GPU

```
x = x.to('cuda')
```

Tensors – Device (GPU)

- Check if your computer has NVIDIA GPU
`torch.cuda.is_available()`
- Multiple GPUs: specify ‘cuda:0’, ‘cuda:1’, ‘cuda:2’, ...

Tensors – Gradient Calculation

```
1 >>> x = torch.tensor([[1., 0.], [-1., 1.]], requires_grad=True)
2 >>> z = x.pow(2).sum()
3 >>> z.backward()
4 >>> x.grad
tensor([[ 2.,  0.],
        [-2.,  2.]])
```

$$\begin{array}{l} \textcircled{1} \\ x = \begin{bmatrix} 1 & 0 \\ -1 & 1 \end{bmatrix} \quad \textcircled{2} \\ z = \sum_i \sum_j x_{i,j}^2 \\ \textcircled{3} \\ \frac{\partial z}{\partial x_{i,j}} = 2x_{i,j} \quad \textcircled{4} \\ \frac{\partial z}{\partial x} = \begin{bmatrix} 2 & 0 \\ -2 & 2 \end{bmatrix} \end{array}$$

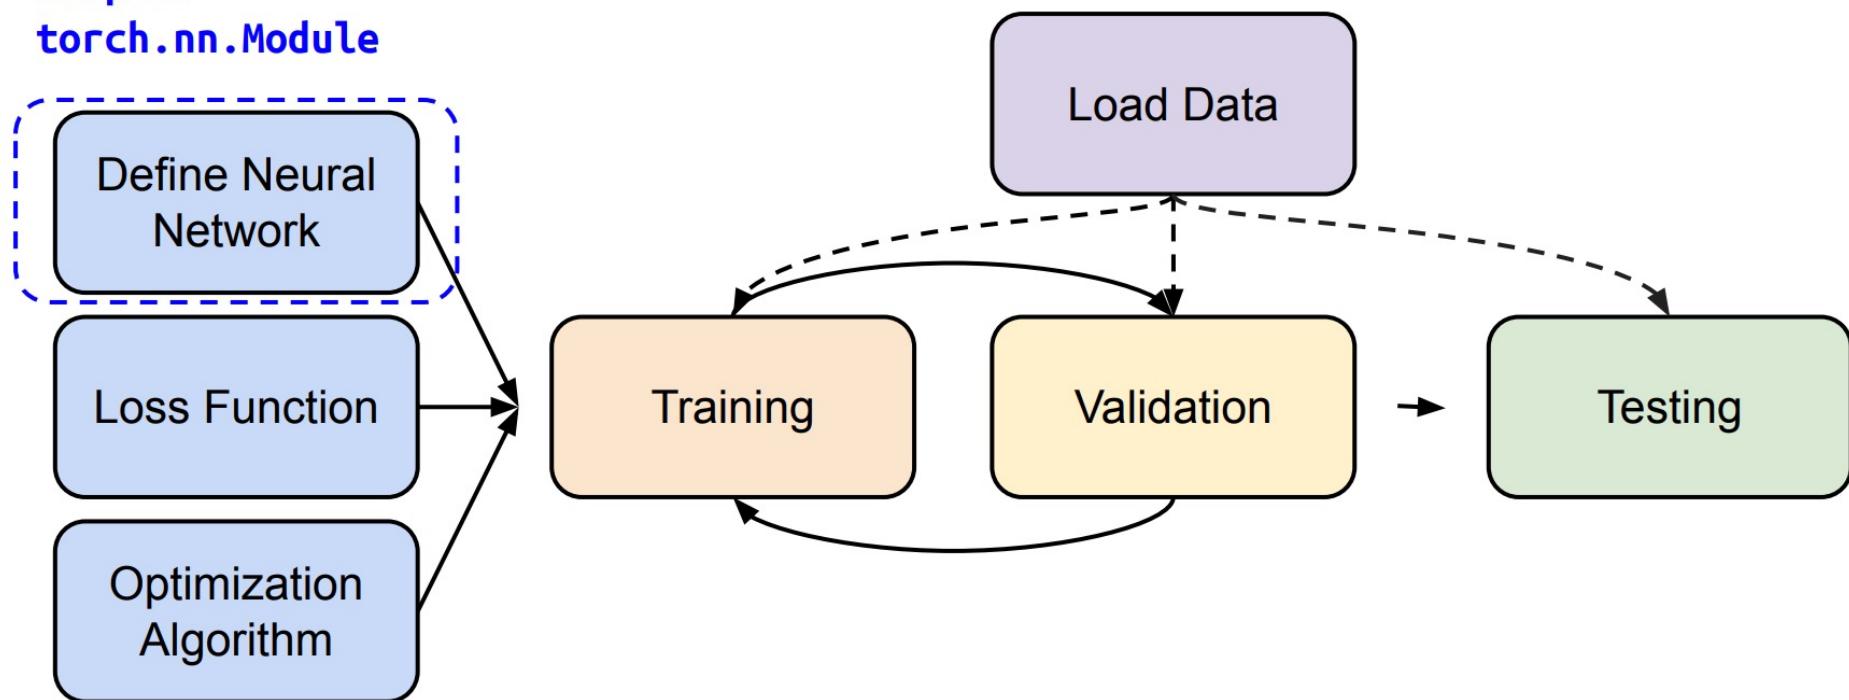
Tensors – Gradient Calculation

```
1 >>> x = torch.tensor([[1., 0.], [-1., 1.]], requires_grad=True)
2 >>> z = x.pow(2).sum()
3 >>> z.backward()
4 >>> x.grad
tensor([[ 2.,  0.],
        [-2.,  2.]])
```

$$\begin{array}{l} \textcircled{1} \\ x = \begin{bmatrix} 1 & 0 \\ -1 & 1 \end{bmatrix} \quad \textcircled{2} \\ z = \sum_i \sum_j x_{i,j}^2 \\ \textcircled{3} \\ \frac{\partial z}{\partial x_{i,j}} = 2x_{i,j} \quad \textcircled{4} \\ \frac{\partial z}{\partial x} = \begin{bmatrix} 2 & 0 \\ -2 & 2 \end{bmatrix} \end{array}$$

Training & Testing NN in PyTorch

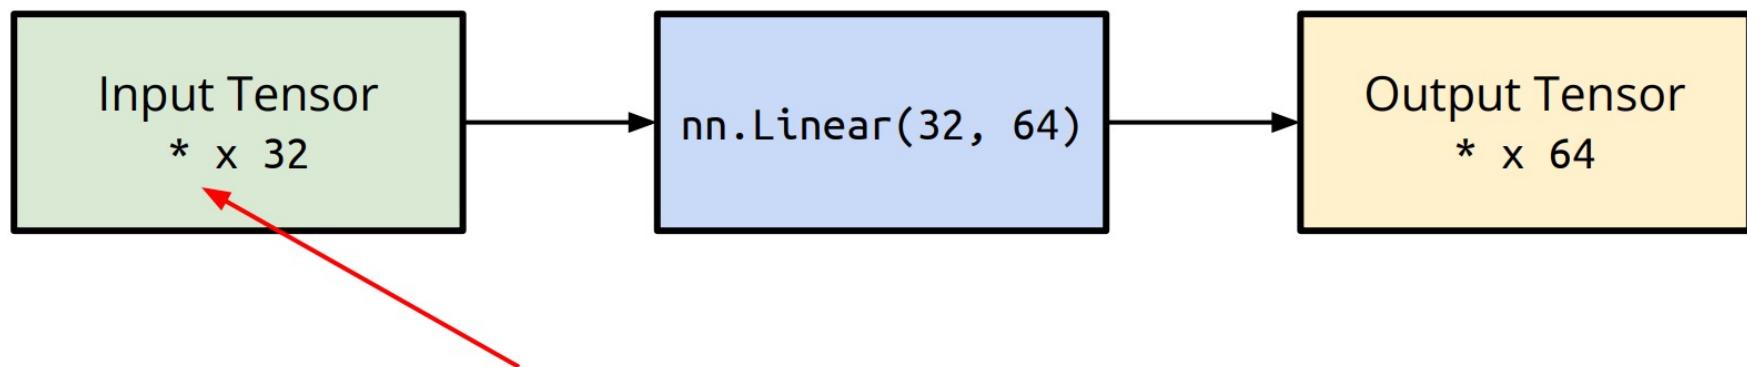
Step 2.
`torch.nn.Module`



Network Layers

- Linear Layer (**Fully-connected** Layer)

```
nn.Linear(in_features, out_features)
```

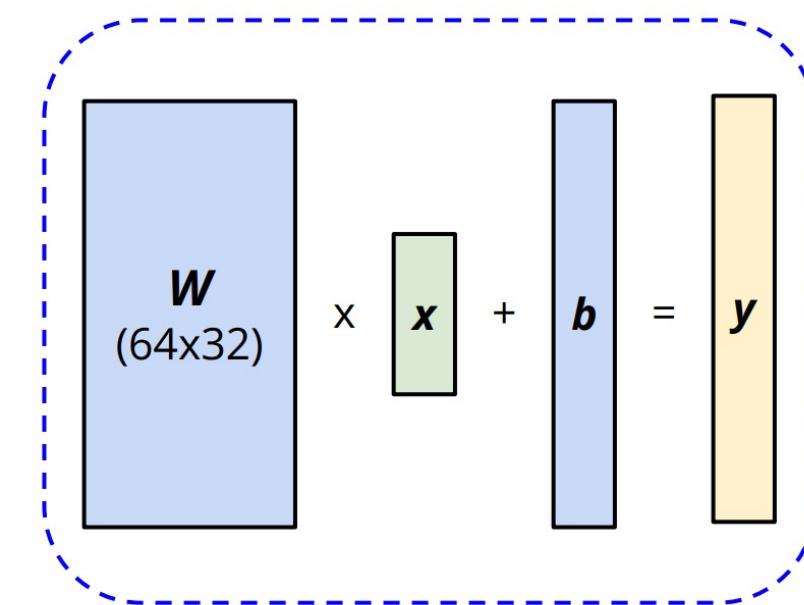
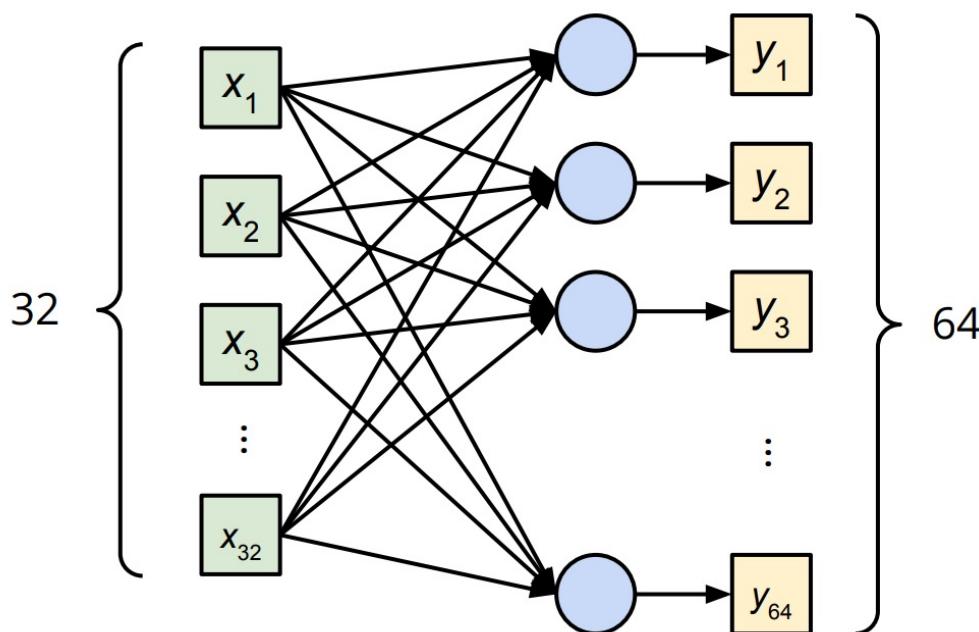


can be any shape (but last dimension must be 32)

e.g. (10, 32), (10, 5, 32), (1, 1, 3, 32), ...

Network Layers

- Linear Layer (**Fully-connected** Layer)



Network Layers

- Linear Layer (**Fully-connected** Layer)

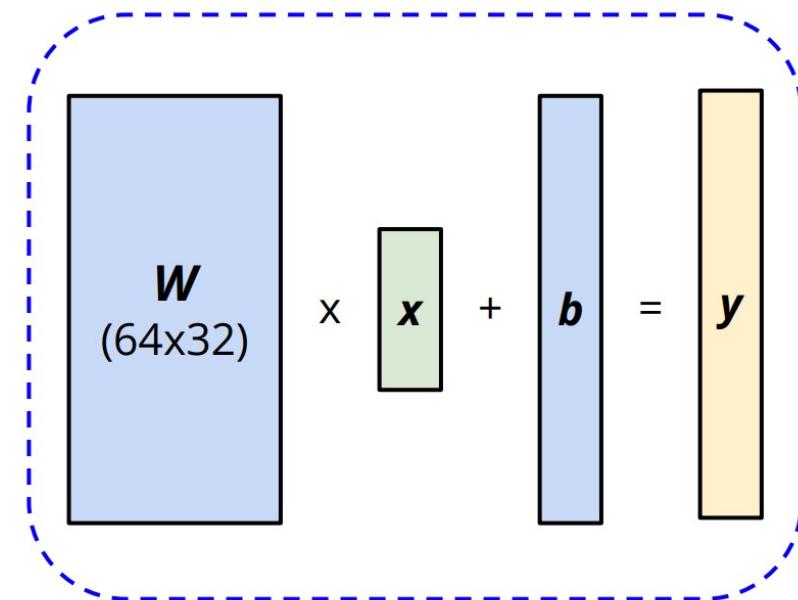
```
>>> layer = torch.nn.Linear(32, 64)
```

```
>>> layer.weight.shape
```

```
torch.Size([64, 32])
```

```
>>> layer.bias.shape
```

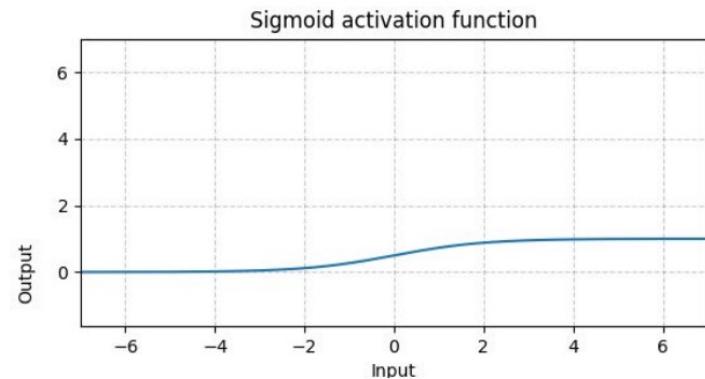
```
torch.Size([64])
```



Non-Linear Activation Functions

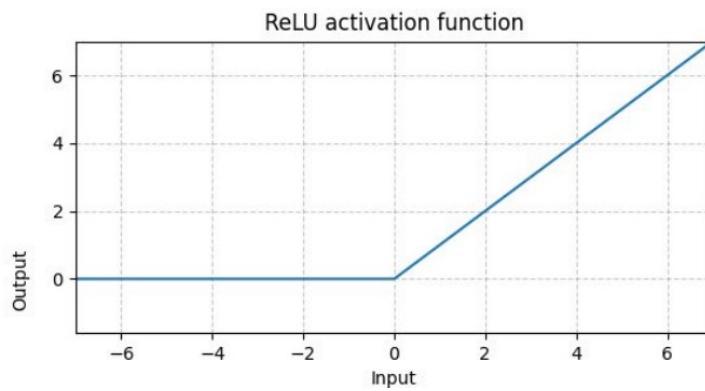
- Sigmoid Activation

`nn.Sigmoid()`



- ReLU Activation

`nn.ReLU()`



Build your own NN

```
import torch.nn as nn
```

```
class MyModel(nn.Module):
    def __init__(self):
        super(MyModel, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(10, 32),
            nn.Sigmoid(),
            nn.Linear(32, 1)
        )
    def forward(self, x):
        return self.net(x)
```



Initialize your model & define layers



Compute output of your NN

Build your own NN

```
import torch.nn as nn

class MyModel(nn.Module):
    def __init__(self):
        super(MyModel, self).__init__()
        self.net = nn.Sequential(
            nn.Linear(10, 32),
            nn.Sigmoid(),
            nn.Linear(32, 1)
        )

    def forward(self, x):
        return self.net(x)
```

```
import torch.nn as nn

class MyModel(nn.Module):
    def __init__(self):
        super(MyModel, self).__init__()
        self.layer1 = nn.Linear(10, 32)
        self.layer2 = nn.Sigmoid(),
        self.layer3 = nn.Linear(32,1)

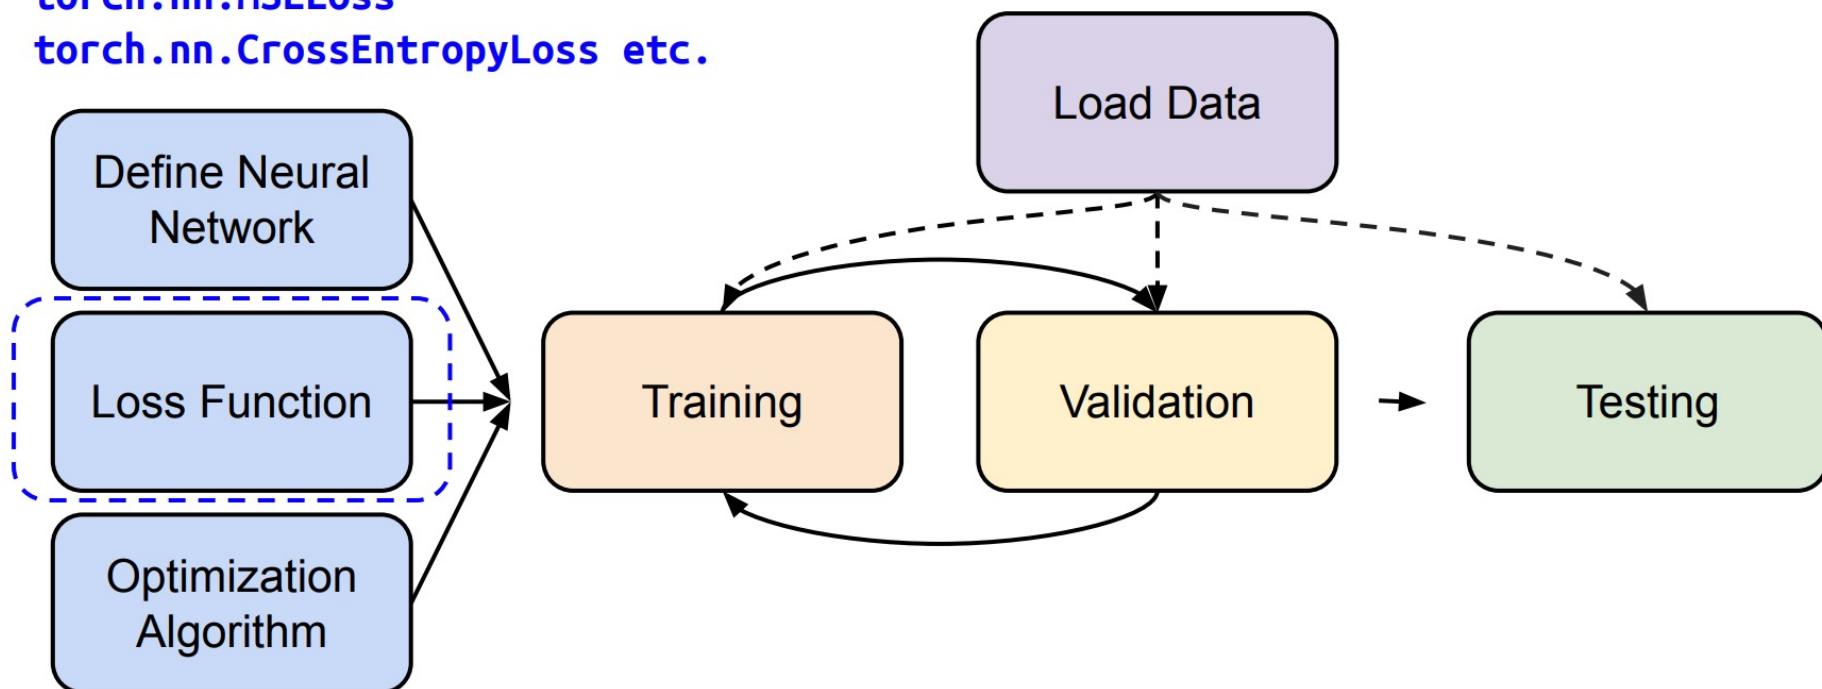
    def forward(self, x):
        out = self.layer1(x)
        out = self.layer2(out)
        out = self.layer3(out)
        return out
```

Training & Testing NN in PyTorch

Step 3.

`torch.nn.MSELoss`

`torch.nn.CrossEntropyLoss` etc.



Loss Functions

- Mean Squared Error (for regression tasks)

```
criterion = nn.MSELoss()
```

- Cross Entropy (for classification tasks)

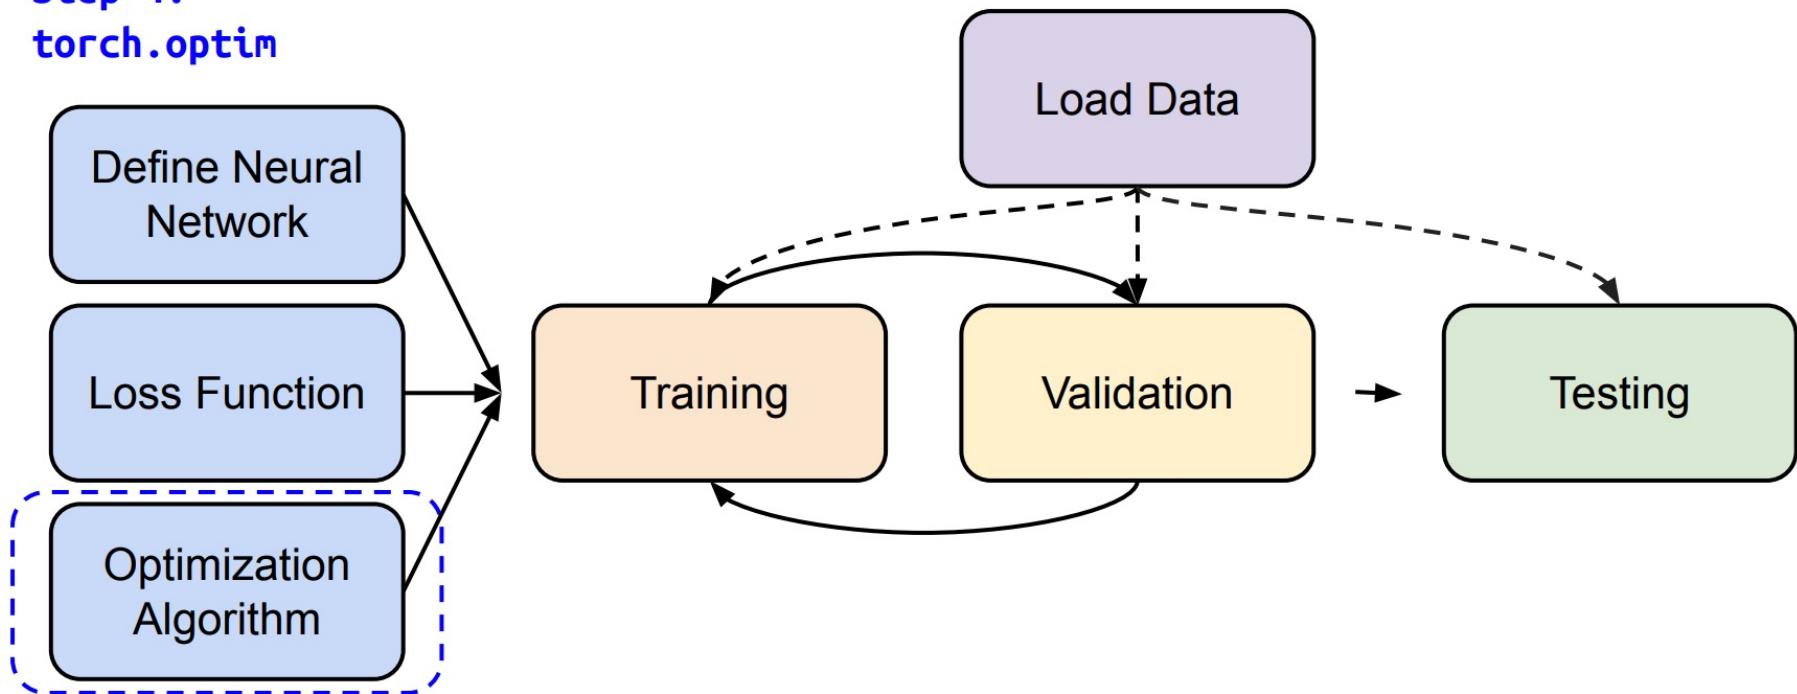
```
criterion = nn.CrossEntropyLoss()
```

- `loss = criterion(model_output, expected_value)`

Training & Testing NN in PyTorch

Step 4.

`torch.optim`



Optimization

- Gradient-based **optimization algorithms** that adjust network parameters to reduce error.
- E.g. Stochastic Gradient Descent (SGD)

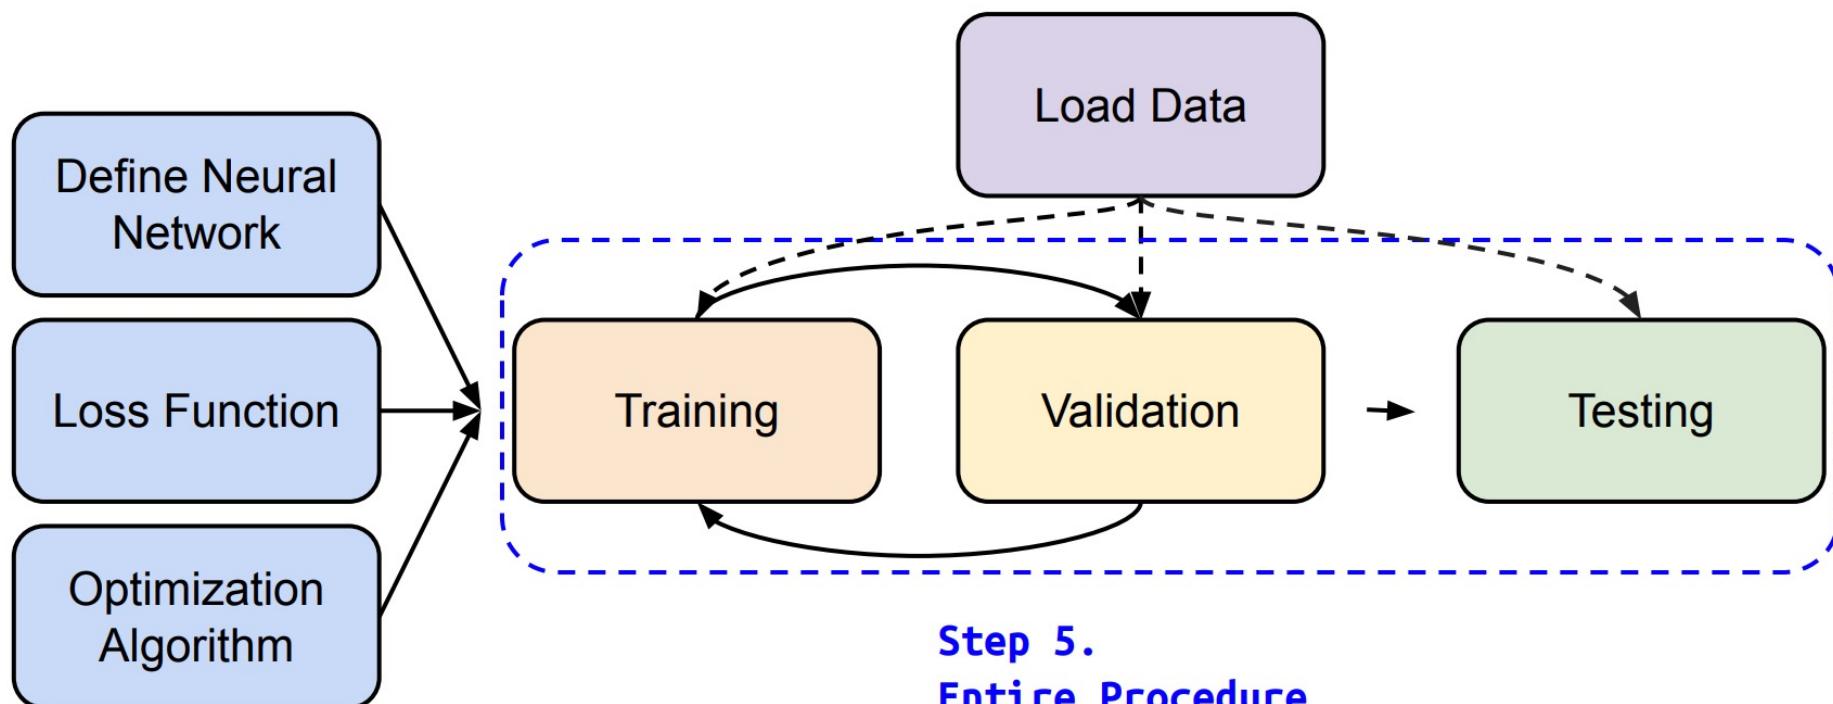
```
torch.optim.SGD(model.parameters(), lr, momentum = 0)
```

Optimization

```
optimizer = torch.optim.SGD(model.parameters(), lr, momentum = 0)
```

- For every batch of data:
 1. Call `optimizer.zero_grad()` to reset gradients of model parameters.
 2. Call `loss.backward()` to backpropagate gradients of prediction loss.
 3. Call `optimizer.step()` to adjust model parameters.

Training & Testing NN in PyTorch



NN Training Setup

```
dataset = MyDataset(file)                                read data via MyDataset  
tr_set = DataLoader(dataset, 16, shuffle=True)          put dataset into Dataloader  
model = MyModel().to(device)                           construct model and move to device (cpu/cuda)  
criterion = nn.MSELoss()                               set loss function  
optimizer = torch.optim.SGD(model.parameters(), 0.1)    set optimizer
```

NN Training Loop

```
for epoch in range(n_epochs):  
    model.train()  
  
    for x, y in tr_set:  
  
        optimizer.zero_grad()  
  
        x, y = x.to(device), y.to(device)  
  
        pred = model(x)  
  
        loss = criterion(pred, y)  
  
        loss.backward()  
  
        optimizer.step()
```

iterate n_epochs
set model to train mode
iterate through the dataloader
set gradient to zero
move data to device (cpu/cuda)
forward pass (compute output)
compute loss
compute gradient (backpropagation)
update model with optimizer

NN Validation Loop

```
model.eval()                                     set model to evaluation mode

total_loss = 0

for x, y in dv_set:                            iterate through the dataloader

    x, y = x.to(device), y.to(device)          move data to device (cpu/cuda)

    with torch.no_grad():                      disable gradient calculation

        pred = model(x)                        forward pass (compute output)

        loss = criterion(pred, y)              compute loss

    total_loss += loss.cpu().item() * len(x)    accumulate loss

avg_loss = total_loss / len(dv_set.dataset)      compute averaged loss
```

NN Testing Loop

```
model.eval()  
preds = []  
  
for x in tt_set:  
    x = x.to(device)  
  
    with torch.no_grad():  
        pred = model(x)  
  
        preds.append(pred.cpu())
```

set model to evaluation mode
iterate through the dataloader
move data to device (cpu/cuda)
disable gradient calculation
forward pass (compute output)
collect prediction

- `model.eval()`
 - Changes behaviour of some model layers, such as dropout and batch normalization.
- `with torch.no_grad()`
 - Prevents calculations from being added into gradient computation graph. Usually used to prevent accidental training on validation/testing data.

Save/Load Trained Models

- Save

```
torch.save(model.state_dict(), path)
```

- Load

```
ckpt = torch.load(path)
```

```
model.load_state_dict(ckpt)
```

More About PyTorch

- torchaudio
 - speech/audio processing
- torchtext
 - natural language processing
- torchvision
 - computer vision
- skorch
 - scikit-learn + pyTorch

More About PyTorch

- torchaudio
 - speech/audio processing
- torchtext
 - natural language processing
- torchvision
 - computer vision
- skorch
 - scikit-learn + pyTorch